

Coming to Dislike Your Opponents: Candidate Evaluations during Presidential Campaigns

Gaurav Sood*

Shanto Iyengar[†]

Abstract

Loud and vitriolic campaigns give partisans many opportunities to reinforce negative views of political opponents. We examine candidate evaluations in the 2000, 2004, and 2008 National Annenberg Election Surveys. Weekly trajectories describe changes over the election year; common autumn windows distinguish late-campaign changes from the longer trajectory. Own-minus-opposing candidate gaps increase across the autumn comparisons, although their size and precision vary. Both improving evaluations of one's own candidate and declining evaluations of the opponent contribute. The balance differs across outcomes and elections. Perceived ideological positions provide context but do not measure the accuracy of respondents' beliefs. These descriptive findings establish patterns that explanations of campaign-period polarization must address. They do not identify the effects of negative advertising or establish greater hostility toward ordinary opposing partisans.

*gsood07@gmail.com

[†]siyengar@stanford.edu

A new chill has descended on inter-partisan relations. Republicans increasingly dislike Democrats, and think most Democrats to be “hypocritical”, and “closed-minded”; for their part, Democrats feel much the same way about Republicans (Iyengar, Sood and Lelkes 2012, see also Iyengar and Westwood 2015; Mason 2013). Not only are partisans increasingly willing to ascribe negative traits to supporters of the opposing party, they are also increasingly socially distant to them. For instance, large proportions of both Democrats and Republicans are “unhappy” at the thought of their relation marrying a supporter of the main opposing party (Iyengar, Sood and Lelkes 2012).

It is against this backdrop of strong affective partisan attachments that the increasingly long, loud, and hostile political campaigns are run. Campaigning in presidential elections now starts in earnest a good two years in advance of the election. On television alone, hundreds of thousands of ads are aired; over a million ads were aired as part of the 2012 general election campaign.¹ Candidates and their surrogates air primarily negative advertisements, and news reports obsess about the most vicious of the negative ads, giving the most perverse of the negative advertisements even greater shelf life (Geer 2010). For instance, during the period preceding the 2004 presidential election, there were more news reports about the notorious Swift Boat Ad questioning Senator Kerry’s military record than on the war in Iraq (Geer 2010).

Partisans process this avalanche of political messages in a manner that favors the party they identify with (Chang 2003; Lodge and Hamill 1986; Rahn 1993). This gives reason to expect warmer feelings for their party’s candidate and greater dislike for the opposing candidate. We examine these expectations using repeated cross sections, separating evaluations of each party’s candidate. We also examine perceived ideological positions, while recognizing that these questions do not directly measure learning.

We assess changes in candidate evaluations using the 2000, 2004, and 2008 National Annenberg Election Surveys. The comparisons describe how evaluations evolve among observed partisans. They do not isolate exposure to campaigning or negative advertising.

We begin with an overview of our theoretical expectations, derived from social identity theory and research showing the reinforcing effects of political campaigns. Next, we describe our data sources, research design, and measurement strategy, present the results, and discuss their implications for ongoing debates over the causes and consequences of mass polarization.

Partisan Processing of Information

Republicans increasingly dislike Democrats, and Democrats feel much the same way about Republicans (Iyengar, Sood and Lelkes 2012). One salient consequence of partisan affect is that it acts as “perceptual screen” (Campbell et al. 1960), encouraging selective exposure and processing of political messages that people receive during campaigns. It is well documented that voters react to political messages not as dispassionate observers, but as biased partisans (Chang 2003; Rahn 1993; Lodge and Hamill 1986; Campbell et al. 1960; Lord, Ross and Lepper 1979). For instance, no matter how flawed the actual performance of candidates during televised debates, partisans are

¹Wesleyan Media Project: <http://mediaproject.wesleyan.edu/2012/11/02/presidential-ad-war-tops-1m-airings/>

quick to declare “their” candidate the winner (Sigelman and Sigelman 1984; Iyengar and Simon 2000). For example, a majority of the Republicans felt that President Ford had out-debated Jimmy Carter despite Ford’s repeated gaffes concerning the autonomy of Eastern Europe in the second debate (Sears and Chaffee 1979).

In case of television ads, which are transparently one-sided, making it relatively easy for partisans to recognize partisanship of the sponsor and react accordingly, partisan filtering is starker still. Partisans accept or reject ads almost instantly based on sponsor of the ad (Iyengar, Jackman and Hahn 2008; McGuire 1969) for the classic discussion of acceptance-rejection factors in the persuasion process, see McGuire 1985. In a series of experiments, Ansolabehere and Iyengar demonstrated that ads proved persuasive only among voters who shared the partisanship of the sponsoring candidate. The reinforcing effect was particularly pronounced among in-partisans with little interest in politics (Ansolabehere and Iyengar 1995, pp. 76–80). Similar experimental findings emerged in studies of the 1992 and 1996 presidential campaigns which confirmed that the reinforcing effect of exposure to campaign ads was especially pronounced among younger voters, who are typically less partisan (Iyengar and Petrocik 2000). These experiments motivate an expectation of partisan reinforcement (for a review see Iyengar and Simon 2000); they do not establish that increasing advertising volume must increase partisan affect in every campaign.

Possible Explanations

The descriptive patterns need not have a single explanation. We consider negative advertising and learning as hypotheses; the available data do not identify their contributions.

A Preference for Negative Advertising

Negative messages may attract more attention than positive appeals (Soroka and McAdams 2012; Meffert et al. 2006), remain more memorable (Lang 1991; Bradley, Angelini and Lee 2007; Chang 2001; Newhagen and Reeves 1991), and lower evaluations of their targets (Gaines and Kuklinski 2011). These findings motivate a hypothesis about candidate evaluations. They do not establish that negative advertising increases the own-opposing gap, or that reactions to candidates extend to parties and their supporters. An attack might also lower evaluations of its sponsor.

Learning

Campaigns may also increase awareness of policy differences between candidates and parties. Earlier work reports modest average gains in campaign knowledge (Geer 2008; Lau, Sigelman and Rovner 2007; Lau and Rovner 2009; Zhao and Chaffee 1995; Ridout et al. 2004; Ansolabehere and Iyengar 1995). The size of that average does not determine how much evaluations change among people who learn consequential facts. Nor does learning necessarily lower evaluations of one’s own candidate: the direction depends on prior beliefs, the positions learned, and the voter’s preferences.

Data and Research Design

Our empirical case rests upon the National Annenberg Election Survey (NAES) rolling cross sections for 2000, 2004, and 2008. The rolling cross sections describe different respondents interviewed at different times. They reveal when and how candidate evaluations differ across the observed campaign periods. Calendar time bundles together advertising, news, debates, nomination decisions, and other events; it does not assign exposure to any one of these influences.

The NAES inputs are archived research workspaces containing the original candidate-question responses, interview dates, and inherited demographic and party-identification variables. They are not complete original survey releases. We reconstruct outcomes from the question responses, treat special nonresponse codes as missing, and retain self-identified Democrats and Republicans. Independent learners enter a sensitivity analysis. The primary estimates are unweighted descriptions of observed respondents. The inherited weights were constructed by the researchers; their full sampling adjustments cannot be certified from these inputs. Both inherited schemes are reported as sensitivities, with unweighted estimates on the same weight-eligible samples available in the replication tables. They do not establish population representativeness.

Measures

Candidate favorability: Respondents rated each presidential candidate on a scale from very unfavorable to very favorable. The original scales have 101 categories in 2000 and 11 categories in 2004 and 2008. We place them on a common numerical range from 0 to 100, retaining their different resolutions. We require both candidate ratings and subtract the opposing-party candidate's rating from the own-party candidate's rating. The difference can range from -100 to 100 ; it is a scale-point difference, not a percentage of people.

Candidate traits: The 2000 items concern caring, honesty, inspiration, and knowledge. The 2004 items concern caring, inspiration, leadership, trustworthiness, knowledge, and recklessness. The 2008 items concern leadership, trustworthiness, experience, and judgment. We map the four ordered 2000 responses to equally spaced values from 0 to 100 and multiply the eleven-category responses in the later years by ten. Recklessness is reversed. Within each respondent, we average only items answered for both candidates, using the same item set for the own- and opposing-candidate averages. Respondents with no paired item are missing. We report complete-item and individual-item sensitivities because question availability and response patterns vary over time. Different trait batteries prevent interpreting levels across years as measurements on an invariant scale. Trait evaluations may express perceived competence and policy-related judgments as well as affect; they do not isolate partisan bias.

Ideological placement: Respondents placed themselves and the candidates on a five-category scale from very conservative to very liberal. We score these categories from 0 to 100 and calculate the absolute distance between self-placement and candidate placement. These are perceived positions. Without an external accuracy criterion, stable average placements cannot establish an absence of learning, and changing distances cannot identify learning as a mediator.

Comparisons and uncertainty

The figures show weekly means during the election year before Election Day. Each election follows a fixed pair of eventual nominees: Bush and Gore in 2000, Bush and Kerry in 2004, and McCain and Obama in 2008. Candidates do not enter or leave these averages as the primary field changes. This retrospective selection describes eventual nominees; it does not describe the entire field of candidates (Appendix C). We use the actual election dates: November 7, 2000; November 2, 2004; and November 4, 2008. Our common autumn comparison is the mean in the final fourteen pre-election days minus the mean on September 1–14. These retrospective windows summarize a common stage of the three campaigns; they are not a preregistered design or a summary of the entire election year. The weekly trajectories remain the primary description of the longer period.

We also standardize the period contrast to the pooled September-through-election-eve respondents with complete outcome and covariate data. Separate linear outcome functions for the two periods include age, gender, Black and Hispanic indicators, education, and party identification, with all slopes allowed to differ by period. Age and the inherited nine-category education score enter linearly. Standardization holds this observed covariate distribution fixed; it does not remove unobserved selection or identify the effect of campaigning. Intervals condition on the empirical standardization target. Complete-case unadjusted contrasts separate the effect of the sample restriction from the effect of adjustment.

Our main NAES intervals use heteroskedasticity-robust covariance clustered by state, with a t reference distribution using the number of states minus one degrees of freedom. The replication tables also give interview-date clustering and respondent-level robust intervals. Neither one-way clustering choice handles every possible dependence across dates and places. Weekly ribbons use respondent-level robust intervals and are descriptive. All intervals are pointwise; the multiple outcomes and specifications should be read together.

We also retain the original comparison of battleground and other states. We fit a linear trend in days toward Election Day, allowing the trend to differ by the historical battleground classification, and cluster by state. Competitiveness, campaign targeting, and the electorate's composition are jointly determined. These slope differences are associations, not estimates of the effect of campaign dosage. The classification uses the original manuscript's state lists. It is not a prospectively assigned treatment.²

²This working draft reports the reconstructed NAES analyses. Rebuilding the original advertising analysis requires the 2000 and 2004 Wisconsin Advertising Project airing records and the geographic crosswalk. Advertising coefficients are not reported until exposure dates, geographic matches, and coverage can be verified.

Results

The weekly candidate gaps describe how the distance between own- and opposing-candidate evaluations evolves over the election year (Figure 1). The 2004 and 2008 favorability trajectories show pronounced widening over the longer period; the 2000 trajectory is flatter and nonmonotonic. In 2008, the common autumn contrast captures only a small part of the longer rise. The gaps should be read alongside the separate own- and opposing-candidate ratings: the same gap can arise from different movements in its two components. These repeated cross sections describe successive samples of partisans, rather than changes within the same people.




Figure 1: Candidate evaluation gaps during the election year. Weekly unweighted own-minus-opposing means among self-identified partisans; ribbons are respondent-level robust 95% intervals. Each outcome requires paired candidate responses. Weeks with fewer than thirty usable observations are omitted. Trait batteries and response availability differ across elections.

Candidate favorability gaps increase across the common autumn windows in all three NAES samples, but the size and precision differ. The increases are 5.7 points in 2000 (95% CI [2.7, 8.7]), 3.0 in 2004 ([0.8, 5.2]), and 2.0 in 2008 ([−0.3, 4.2]). The unadjusted 2008 interval includes zero. Standardization gives positive estimates in each year, conditional on the observed covariates and modeling choices (Table 1). The numerical similarity between adjusted and unadjusted estimates does not establish that selection is absent.

The weekly trajectories show why the window matters (Figure 2). A change from early in the election year includes developments before the general-election period, including candidate

Table 1: Autumn changes in candidate evaluation gaps

Year	Measure	Early	Late	Change [95% CI]	Adjusted change
2000	Favorability gap	39.3	45.0	5.7 [2.7, 8.7]	5.8 [3.0, 8.5]
2000	Trait gap	25.7	31.6	6.0 [4.2, 7.7]	5.9 [4.4, 7.5]
2004	Favorability gap	54.2	57.1	3.0 [0.8, 5.2]	3.2 [1.0, 5.4]
2004	Trait gap	41.6	44.9	3.3 [0.8, 5.8]	3.8 [1.4, 6.3]
2008	Favorability gap	43.6	45.5	2.0 [-0.3, 4.2]	2.4 [0.3, 4.4]
2008	Trait gap	36.6	39.3	2.7 [0.6, 4.7]	2.9 [1.1, 4.6]

Early: September 1–14. Late: final fourteen days before Election Day. All quantities are scale points. Gaps are own minus opposing candidate; changes are late minus early. Primary estimates are unweighted. Adjusted contrasts standardize to complete-case autumn respondents. Both interval columns use state clustering. Samples and alternative inference are in the replication tables.

selection and changing familiarity. It should not be read as the effect of a fixed dose of campaigning. Weekly observations also describe changing sets of respondents, not individual trajectories.




Figure 2: Own- and opposing-candidate favorability during the election year. Weekly unweighted means among self-identified partisans with both ratings; ribbons are respondent-level robust 95% intervals. Weeks with fewer than thirty paired observations are omitted.

Disaggregating the gap shows increases in own-candidate favorability and decreases in opposing-candidate favorability in each autumn comparison (Table 2). The opposing-candidate declines are numerically larger, but the interval for the sum of the two changes includes zero in each year. Thus these comparisons do not establish that declining evaluations of opponents consistently dominate improving evaluations of one’s own candidate.

Trait gaps also increase (Table 1; Figure 3). Their components differ across elections. In 2004, the increase in own-candidate trait ratings is larger than the decline in opposing-candidate ratings. Complete-item and inherited-weight sensitivities widen the uncertainty around some trait comparisons, including intervals that cross zero (Appendix A). Question availability and the par-

ticular traits measured matter to the conclusion.

Table 2: Own- and opposing-candidate components of autumn changes

Year	Measure	Candidate	Change [95% CI]	<i>N</i>
2000	Favorability	Own candidate	2.6 [1.2, 4.1]	4767
2000	Favorability	Opposing candidate	-3.0 [-5.2, -0.9]	4767
2000	Traits	Own candidate	1.0 [-0.1, 2.1]	4889
2000	Traits	Opposing candidate	-5.0 [-6.1, -3.8]	4889
2004	Favorability	Own candidate	1.2 [0.2, 2.2]	6159
2004	Favorability	Opposing candidate	-1.8 [-3.4, -0.3]	6159
2004	Traits	Own candidate	2.6 [1.3, 3.8]	4031
2004	Traits	Opposing candidate	-0.7 [-2.2, 0.8]	4031
2008	Favorability	Own candidate	0.4 [-0.9, 1.8]	4461
2008	Favorability	Opposing candidate	-1.5 [-2.9, -0.1]	4461
2008	Traits	Own candidate	0.5 [-0.7, 1.8]	4450
2008	Traits	Opposing candidate	-2.1 [-3.4, -0.9]	4450

Unweighted late-minus-early changes, with state-clustered 95% intervals. *N* counts respondents in the two comparison periods. Both components use the same outcome-specific sample.




Figure 3: Candidate trait evaluations. Each respondent's means use traits answered for both candidates. Batteries differ by year. Weekly unweighted means and respondent-level robust 95% intervals; weeks with fewer than thirty usable observations are omitted.

The battleground comparisons show faster growth in trait gaps in each election, with point-wise intervals excluding zero (Table 3). Favorability gaps do not show the same consistency: only the 2008 slope difference has an interval excluding zero. These comparisons therefore provide mixed descriptive evidence for the expectation that candidate evaluations diverge more rapidly in competitive states. Differences in targeting, state electorates, and other campaign developments remain alternative explanations.

Finally, perceived ideological positions and self-candidate distances provide descriptive con-

Table 3: Differences in trends between battleground and other states

Year	Measure	Slope difference [95% CI]	<i>N</i>
2000	Favorability gap	-0.5 [-1.5, 0.5]	28164
2000	Trait gap	1.4 [0.4, 2.4]	25589
2004	Favorability gap	1.1 [-0.2, 2.4]	40678
2004	Trait gap	1.1 [0.1, 2.2]	26820
2008	Favorability gap	1.3 [0.3, 2.2]	34048
2008	Trait gap	1.9 [0.7, 3.1]	24359

Battleground minus other-state linear slopes, in scale points per 100 days toward Election Day. Unweighted election-year pre-election observations; state-clustered 95% intervals. Historical state lists are held fixed within each election. These are descriptive slope differences, not effects of campaign exposure.

text (Appendix B). They cannot adjudicate between learning, projection, and identity-based evaluation. Average positions can stay stable while some respondents learn and others change in the opposite direction. Item nonresponse and the coarse response scale further limit what a stable mean can reveal.

Discussion

Our research connects to ongoing research showing that partisanship is a salient basis for group identity (Iyengar, Sood and Lelkes 2012; Freeze 2012). Across three presidential elections, the NAES data show growing candidate evaluation gaps during the common autumn periods, with differences in magnitude and precision. The weekly trajectories place those contrasts within the longer election year. The findings characterize a phenomenon worth explaining even when these data cannot identify its causes.

The components of the gap matter. Increasing warmth toward one’s own candidate and declining evaluations of the opposing candidate both contribute, and their relative importance differs across elections and measures. A widening gap alone cannot establish that hostility is the dominant change. Nor do unfavorable evaluations of named candidates establish hostility toward ordinary supporters of the opposing party. Candidate evaluations can respond to perceived competence, policy disagreement, identity, and other considerations.

Partisan processing of information, strengthened attachments, and more salient identities remain plausible explanations for increasing candidate gaps. So do changing beliefs about candidates’ competence and policy positions. Our perceived-position questions do not measure learning directly. Stable average placements could coexist with individual learning, offsetting changes, or changes too small for the response categories to capture. The data therefore leave learning’s contribution open.

The comparisons across battleground and other states add geographic variation to the description. They do not isolate the effects of greater campaign activity. Similarly, a relationship between advertising volume and evaluations would require credible exposure measurement before it could be interpreted even as an association. A causal interpretation would further require an argument that separates advertising from campaign targeting and the other differences between places.

The broader 2008 candidate field shows why selection of eventual nominees matters. Opposing-party evaluations of Clinton rise over the observed year, while those of Obama and McCain decline (Appendix C). Several other candidates cease to be measured early. Their unobserved later ratings cannot be supplied by averaging the candidates who remain. The nominee trajectories therefore do not establish a common path for all presidential contenders.

The results concern observed respondents in three historical presidential campaigns. Changes in questionnaire availability, response patterns, and candidate familiarity affect which comparisons are informative. Covariate adjustment and alternative scoring choices probe particular sources of sensitivity, while leaving unmeasured differences unresolved. The descriptive contribution is to show the timing, components, and variation of candidate polarization clearly enough to guide more specific explanatory tests.

References

- Ansolabehere, Stephen and Shanto Iyengar. 1995. *Going negative*. Simon and Schuster.
- Bradley, Samuel D, James R Angelini and Sungkyoung Lee. 2007. "Psychophysiological and memory effects of negative political ads: Aversive, arousing, and well remembered." *Journal of Advertising* 36(4):115–127.
- Campbell, Angus, Philip E Converse, Warren E Miller and Donald E Stokes. 1960. *The American Voter*. New York: Wiley.
- Chang, Chingching. 2001. "The impacts of emotion elicited by print political advertising on candidate evaluation." *Media psychology* 3(2):91–118.
- Chang, Chingching. 2003. "Party Bias in Political-Advertising Processing—Results from an Experiment Involving the 1998 Taipei Mayoral Election." *Journal of Advertising* 32(2):55–67.
- Freeze, Melanie Sue. 2012. "Rallying Around the Party: A Theory of Party Identity Linkage." Working Paper.
- Gaines, Brian J and James H Kuklinski. 2011. "Experimental Estimation of Heterogeneous Treatment Effects Related to Self-Selection." *American Journal of Political Science* 55(3):724–736.
- Geer, John G. 2008. *In defense of negativity: Attack ads in presidential campaigns*. University of Chicago Press.
- Geer, John G. 2010. *Fanning the Flames: The News Media's Role in the Rise of Negativity in Presidential Campaigns*. Technical report Discussion Paper, Joan Shorenstein Center on the Press, Politics, and Public Policy, Harvard University.
- Iyengar, Shanto and Adam F Simon. 2000. "New perspectives and evidence on political communication and campaign effects." *Annual review of psychology* 51(1):149–169.
- Iyengar, Shanto, Gaurav Sood and Yphtach Lelkes. 2012. "Affect, Not Ideology A Social Identity Perspective on Polarization." *Public opinion quarterly* 76(3):405–431.
- Iyengar, Shanto and John R Petrocik. 2000. "'Basic Rule' Voting: Impact of Campaigns on Party-and Approval-Based Voting." *Crowded airwaves: Campaign advertising in elections* pp. 113–148.
- Iyengar, Shanto and Sean J. Westwood. 2015. "Fear and Loathing Across Party Lines: New Evidence on Group Polarization." *American Journal of Political Science* 59(3):690–707.
- Iyengar, Shanto, Simon Jackman and Kyu Hahn. 2008. Polarization in less than thirty seconds: continuous monitoring of voter response to campaign advertising. In *annual meetings of the Midwest Political Science Association, Chicago, IL*. Citeseer.
- Lang, Annie. 1991. "Emotion, formal features, and memory for televised political advertisements." *Television and political advertising* 1:221–243.
- Lau, Richard R and Ivy Brown Rovner. 2009. "Negative campaigning." *Annual review of political science* 12:285–306.
- Lau, Richard R, Lee Sigelman and Ivy Brown Rovner. 2007. "The effects of negative political campaigns: a meta-analytic reassessment." *Journal of Politics* 69(4):1176–1209.
- Lodge, Milton and Ruth Hamill. 1986. "A partisan schema for political information processing." *The American Political Science Review* pp. 505–520.
- Lord, Charles G, Lee Ross and Mark R Lepper. 1979. "Biased assimilation and attitude polarization: The effects of prior theories on subsequently considered evidence." *Journal of Personality and Social Psychology* 37(11):2098.
- Mason, Lilliana. 2013. "The Rise of Uncivil Agreement: Issue versus Behavioral Polarization in the American Electorate." *American Behavioral Scientist* 57(1):140–159.
- McGuire, William J. 1969. "The nature of attitudes and attitude change." *The Handbook of Social Psychology* 3:136–314.
- McGuire, William J. 1985. Attitudes and Attitude Change. In *Handbook of Social Psychology*, ed. Gardner Lindzey and Elliot Aronson. Knopf.
- Meffert, Michael F, Sungeun Chung, Amber J Joiner, Leah Waks and Jennifer Garst. 2006. "The effects of

- negativity and motivated information processing during a political campaign." *Journal of Communication* 56(1):27–51.
- Newhagen, John E and Byron Reeves. 1991. "Emotion and memory responses for negative political advertising: A study of television commercials used in the 1988 presidential election." *Television and political advertising* 1:197–220.
- Rahn, Wendy M. 1993. "The role of partisan stereotypes in information processing about political candidates." *American Journal of Political Science* pp. 472–496.
- Ridout, Travis N, Dhavan V Shah, Kenneth M Goldstein and Michael M Franz. 2004. "Evaluating measures of campaign advertising exposure on political learning." *Political Behavior* 26(3):201–225.
- Sears, David O and Steven H Chaffee. 1979. Uses and effects of the 1976 debates: An overview of empirical studies. In *The Great Debates: Carter vs. Ford, 1976*, ed. Sidney Kraus. Indiana University Press pp. 223–261.
- Sigelman, Lee and Carol K Sigelman. 1984. "Judgments of the Carter-Reagan debate: The eyes of the beholders." *Public Opinion Quarterly* 48(3):624–628.
- Soroka, Stuart and Stephen McAdams. 2012. News, Politics, and Negativity. Technical report CIRANO.
- Zhao, Xinshu and Steven H Chaffee. 1995. "Campaign advertisements versus television news as sources of political issue information." *Public Opinion Quarterly* 59(1):41–65.

A Sensitivity and response availability

Table 4 varies weights, trait completeness, and the partisan sample. The inherited weights are researcher-created and are not treated as verified study weights. The full machine-readable tables additionally report individual traits, matched-sample unweighted comparisons, demographic composition, floor and ceiling frequencies, battleground trend differences, and alternative variance estimators. These are sensitivity analyses of descriptive contrasts; they do not create a counterfactual campaign exposure.

Table 4: Alternative NAES autumn specifications

Year	Measure	Specification	Change [95% CI]	<i>N</i>
2000	Favorability gap	Inherited weight 1	5.9 [2.3, 9.5]	4672
2000	Trait gap	Inherited weight 1	5.6 [3.7, 7.5]	4788
2000	Favorability gap	Inherited weight 2	6.6 [2.2, 10.9]	4672
2000	Trait gap	Inherited weight 2	5.9 [3.6, 8.2]	4788
2000	Trait gap	All trait pairs	5.9 [3.9, 8.0]	4435
2004	Favorability gap	Inherited weight 1	2.9 [0.0, 5.8]	6030
2004	Trait gap	Inherited weight 1	3.1 [-0.4, 6.7]	3950
2004	Favorability gap	Inherited weight 2	2.6 [-0.2, 5.4]	6030
2004	Trait gap	Inherited weight 2	2.8 [-0.5, 6.2]	3950
2004	Trait gap	All trait pairs	2.6 [-0.1, 5.2]	3657
2008	Favorability gap	Inherited weight 1	1.1 [-1.7, 3.8]	4258
2008	Trait gap	Inherited weight 1	2.5 [-0.1, 5.1]	4247
2008	Favorability gap	Inherited weight 2	1.2 [-1.8, 4.2]	4258
2008	Trait gap	Inherited weight 2	2.7 [-0.2, 5.6]	4247
2008	Trait gap	All trait pairs	2.8 [0.6, 5.0]	4206
2000	Favorability gap	Including leaners	4.9 [3.3, 6.6]	7092
2000	Trait gap	Including leaners	4.9 [3.8, 6.1]	7289
2004	Favorability gap	Including leaners	3.6 [1.8, 5.5]	8450
2004	Trait gap	Including leaners	4.0 [2.2, 5.9]	5478
2008	Favorability gap	Including leaners	1.8 [0.0, 3.6]	6195
2008	Trait gap	Including leaners	2.5 [0.8, 4.2]	6174

State-clustered 95% intervals. Samples differ with weight and item availability. All trait pairs requires responses for every trait for both candidates. Including leaners uses the inherited seven-category identification measure.




Figure 4: Weekly response availability among self-identified NAES partisans. The denominator includes respondents without usable outcome responses. Availability combines questionnaire fielding and respondent nonresponse.

B Perceived ideology




Figure 5: Perceived ideological placement of candidates. The scale runs from conservative (0) to liberal (100); rows separate Democratic and Republican respondents. Each series uses its available responses. Weekly unweighted means with respondent-level robust 95% intervals, omitting weeks with fewer than thirty usable responses.




Figure 6: Absolute distance between self-placement and perceived candidate placement. Each series requires both the respondent's and candidate's placement. A distance is not a measure of accuracy. Weekly unweighted means with respondent-level robust 95% intervals, omitting weeks with fewer than thirty usable responses.

C The broader candidate field in 2008

The manuscript holds candidate identities fixed within each election. The figure below checks whether the 2008 nominee trajectories characterize other observed presidential contenders. Each candidate's available responses are retained separately. The nine candidates are those with verified presidential-contender favorability items in the archived workspace; this is not a census of all entrants. Biden and Palin ratings begin during the vice-presidential campaign and are excluded. Question availability does not define the active primary field: Clinton continues to be rated after her primary campaign, whereas Paul's item is asked only briefly in January. The official NAES08 phone codebook documents the fielding schedules.³

³<https://cdn.annenbergpublicpolicycenter.org/Downloads/NAES/PhoneSurvey/NAES08-Phone-Codebook.pdf>, variable directory, pp. 37–43.




Figure 7: Favorability of individual 2008 presidential contenders. Weekly unweighted means among self-identified partisans, separated by whether the respondent shares the candidate's party. Candidate-specific available responses; pointwise respondent-level robust 95% intervals. Weeks with fewer than thirty ratings in a group are omitted. Blank periods are unobserved or insufficiently measured, not zero ratings or verified withdrawal dates.